

# DS Automation Assignment

Using our prepared churn data from week 2:

- use pycaret to find an ML algorithm that performs best on the data
  - Choose a metric you think is best to use for finding the best model; by default, it is accuracy but it could be AUC, precision, recall, etc. The week 3 FTE has some information on these different metrics.
- save the model to disk
- create a Python script/file/module with a function that takes a pandas dataframe as an input and returns the probability of churn for each row in the dataframe
  - your Python file/function should print out the predictions for new data (new\_churn\_data.csv)
  - the true values for the new data are [1, 0, 0, 1, 0] if you're interested
- test your Python module and function with the new data, new\_churn\_data.csv
- write a short summary of the process and results at the end of this notebook
- upload this Jupyter Notebook and Python file to a Github repository, and turn in a link to the repository in the week 5 assignment dropbox

*Optional challenges:*

- return the probability of churn for each new prediction, and the percentile where that prediction is in the distribution of probability predictions from the training dataset (e.g. a high probability of churn like 0.78 might be at the 90th percentile)
- use other autoML packages, such as TPOT, H2O, MLBox, etc, and compare performance and features with pycaret
- create a class in your Python module to hold the functions that you created
- accept user input to specify a file using a tool such as Python's `input()` function, the `click` package for command-line arguments, or a GUI
- Use the unmodified churn data (new\_unmodified\_churn\_data.csv) in your Python script. This will require adding the same preprocessing steps from week 2 since this data is like the original unmodified dataset from week 1.

## Load Data

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv('prepped_churn_data.csv', index_col='customerID')
#df = pd.read_csv('new_churn_data.csv', index_col='customerID')
df
```

Out[2]:

	tenure	PhoneService	Contract	PaymentMethod	MonthlyCharges	TotalCharges	Churn
customerID							
7590-VHVEG	1	0	0	1	29.85	29.85	0
5575-GNVDE	34	1	1	0	56.95	1889.50	0
3668-QPYBK	2	1	0	0	53.85	108.15	1
7795-CFOCW	45	0	1	2	42.30	1840.75	0
9237-HQITU	2	1	0	1	70.70	151.65	1
...	...	...	...	...	...	...	...
6840-RESVB	24	1	1	0	84.80	1990.50	0
2234-XADUH	72	1	1	3	103.20	7362.90	0
4801-JZAZL	11	0	0	1	29.60	346.45	0
8361-LTMKD	4	1	0	0	74.40	306.60	1
3186-AJIEK	66	1	2	2	105.65	6844.50	0

7032 rows × 8 columns



## AutoML with pycaret

```
In [3]: from pycaret.classification import *
```

```
In [4]: automl = setup(data = df, target = 'Churn', fold_shuffle=True)
```

	Description	Value
0	session_id	8860
1	Target	Churn
2	Target Type	Binary
3	Label Encoded	0: 0, 1: 1
4	Original Data	(7032, 8)
5	Missing Values	False
6	Numeric Features	4
7	Categorical Features	3
8	Ordinal Features	False
9	High Cardinality Features	False
10	High Cardinality Method	None
11	Transformed Train Set	(4922, 12)
12	Transformed Test Set	(2110, 12)
13	Shuffle Train-Test	True
14	Stratify Train-Test	False
15	Fold Generator	StratifiedKFold
16	Fold Number	10
17	CPU Jobs	-1
18	Use GPU	False
19	Log Experiment	False
20	Experiment Name	clf-default-name
21	USI	a8f7
22	Imputation Type	simple
23	Iterative Imputation Iteration	None
24	Numeric Imputer	mean
25	Iterative Imputation Numeric Model	None
26	Categorical Imputer	constant
27	Iterative Imputation Categorical Model	None
28	Unknown Categoricals Handling	least_frequent
29	Normalize	False
30	Normalize Method	None
31	Transformation	False
32	Transformation Method	None

	Description	Value
33	PCA	False
34	PCA Method	None
35	PCA Components	None
36	Ignore Low Variance	False
37	Combine Rare Levels	False
38	Rare Level Threshold	None
39	Numeric Binning	False
40	Remove Outliers	False
41	Outliers Threshold	None
42	Remove Multicollinearity	False
43	Multicollinearity Threshold	None
44	Clustering	False
45	Clustering Iteration	None
46	Polynomial Features	False
47	Polynomial Degree	None
48	Trigonometry Features	False
49	Polynomial Threshold	None
50	Group Features	False
51	Feature Selection	False
52	Features Selection Threshold	None
53	Feature Interaction	False
54	Feature Ratio	False
55	Interaction Threshold	None
56	Fix Imbalance	False
57	Fix Imbalance Method	SMOTE

In [5]: `automl[6]`

Out[5]: -1

In [6]: `best_model = compare_models()`

	Model	Accuracy	AUC	Recall	Prec.	F1	Kappa	MCC	TT (Sec)
<b>lda</b>	Linear Discriminant Analysis	0.7901	0.8325	0.5045	0.6445	0.5650	0.4295	0.4357	0.0030
<b>lr</b>	Logistic Regression	0.7891	0.8380	0.4940	0.6454	0.5589	0.4235	0.4306	0.2280
<b>ridge</b>	Ridge Classifier	0.7885	0.0000	0.4571	0.6584	0.5386	0.4072	0.4192	0.0030
<b>gbc</b>	Gradient Boosting Classifier	0.7859	0.8345	0.4910	0.6341	0.5529	0.4152	0.4214	0.0410
<b>ada</b>	Ada Boost Classifier	0.7816	0.8311	0.4789	0.6258	0.5420	0.4021	0.4086	0.0160
<b>catboost</b>	CatBoost Classifier	0.7814	0.8297	0.4940	0.6203	0.5496	0.4078	0.4126	0.2710
<b>lightgbm</b>	Light Gradient Boosting Machine	0.7735	0.8197	0.4947	0.5985	0.5413	0.3928	0.3962	0.0100
<b>xgboost</b>	Extreme Gradient Boosting	0.7733	0.8083	0.4857	0.6000	0.5366	0.3887	0.3927	0.0530
<b>rf</b>	Random Forest Classifier	0.7643	0.7853	0.4759	0.5784	0.5216	0.3674	0.3708	0.0350
<b>knn</b>	K Neighbors Classifier	0.7574	0.7336	0.4308	0.5664	0.4889	0.3340	0.3395	0.1250
<b>et</b>	Extra Trees Classifier	0.7475	0.7581	0.4677	0.5407	0.5005	0.3329	0.3351	0.0350
<b>dt</b>	Decision Tree Classifier	0.7170	0.6512	0.4925	0.4786	0.4849	0.2901	0.2904	0.0030
<b>svm</b>	SVM - Linear Kernel	0.7109	0.0000	0.4293	0.5646	0.4199	0.2574	0.2996	0.0040
<b>nb</b>	Naive Bayes	0.7064	0.8159	0.8075	0.4749	0.5979	0.3904	0.4248	0.0940
<b>qda</b>	Quadratic Discriminant Analysis	0.5645	0.5271	0.4459	0.2066	0.2616	0.0385	0.0525	0.0030

In [7]: `best_model`

Out[7]: `LinearDiscriminantAnalysis(n_components=None, priors=None, shrinkage=None, solver='svd', store_covariance=False, tol=0.0001)`

In [8]: `df.iloc[-2:-1].shape`

Out[8]: `(1, 8)`

In [9]: `predict_model(best_model, df.iloc[-2:-1])`

Out[9]:

	tenure	PhoneService	Contract	PaymentMethod	MonthlyCharges	TotalCharges	Churn
customerID							
8361-LTMKD	4	1	0	0	74.4	306.6	1

## Saving and Loading Model

In [10]: `save_model(best_model, 'LDA')`

Transformation Pipeline and Model Successfully Saved

Out[10]:

```
(Pipeline(memory=None,
          steps=[('dtypes',
                  DataTypes_Auto_infer(categorical_features=[],
                                       display_types=True, features_todrop=[],
                                       id_columns=[],
                                       ml_usecase='classification',
                                       numerical_features=[], target='Churn',
                                       time_features=[])),
                  ('imputer',
                   Simple_Imputer(categorical_strategy='not_available',
                                  fill_value_categorical=None,
                                  fill_value_numerical=None,
                                  numeric_strate...
                  ('dummy', Dummify(target='Churn')),
                  ('fix_perfect', Remove_100(target='Churn')),
                  ('clean_names', Clean_Colum_Names()),
                  ('feature_select', 'passthrough'), ('fix_multi', 'passthrough'),
                  ('dfs', 'passthrough'), ('pca', 'passthrough'),
                  ['trained_model',
                   LinearDiscriminantAnalysis(n_components=None, priors=None,
                                              shrinkage=None, solver='svd',
                                              store_covariance=False,
                                              tol=0.0001)]],
          verbose=False),
        'LDA.pkl')
```

In [11]: `import pickle`

```
with open('LDA_model.pk', 'wb') as f:
    pickle.dump(best_model, f)
```

In [12]: `with open('LDA_model.pk', 'rb') as f:`  
`loaded_model = pickle.load(f)`

In [13]: `new_data = df.iloc[-2:-1].copy()`  
`#new_data.drop('MonthlyCharges_tenure_ratio', axis=1, inplace=True)`  
`new_data['dummy1']=0`  
`new_data['dummy2']=0`  
`new_data['dummy3']=0`  
`new_data['dummy4']=0`  
`new_data`

Out[13]:

	tenure	PhoneService	Contract	PaymentMethod	MonthlyCharges	TotalCharges	Churn
customerID							
8361-LTMKD	4	1	0	0	74.4	306.6	1

In [14]: `loaded_model.predict(new_data)`

Out[14]: `array([1])`

```
In [15]: loaded_lda = load_model('LDA')
```

Transformation Pipeline and Model Successfully Loaded

```
In [16]: predict_model(loaded_lda, new_data)
```

```
Out[16]:
```

	tenure	PhoneService	Contract	PaymentMethod	MonthlyCharges	TotalCharges	Churn
customerID							
8361-LTMKD	4	1	0	0	74.4	306.6	1

## Python Predictions Module

```
In [17]: from IPython.display import Code
```

```
Code('predict_churn.py')
```

```
Out[17]: import pandas as pd
from pycaret.classification import predict_model, load_model
```

```
def load_data(filepath):
    """
    Loads churn data into a DataFrame from a string filepath.
    """
    df = pd.read_csv(filepath, index_col='customerID')
    return df

def make_predictions(df):
    """
    Uses the pycaret best model to make predictions on data in the df dataframe.
    """
    model = load_model('LDA')
    predictions = predict_model(model, data=df)
    predictions.rename({'Label': 'Churn_prediction'}, axis=1, inplace=True)
    predictions['Churn_prediction'].replace({1: 'Churn', 0: 'No churn'},
                                             inplace=True)
    return predictions['Churn_prediction']

if __name__ == "__main__":
    df = load_data('new_churn_data.csv')
    predictions = make_predictions(df)
    print('predictions:')
    print(predictions)
```

```
In [18]: %run predict_churn.py
```

Transformation Pipeline and Model Successfully Loaded

predictions:

customerID

9305-CKSKC           Churn

1452-KNGVK           Churn

6723-OKKJM       No churn

7832-POPKP           Churn

6348-TACGU           Churn

Name: Churn\_prediction, dtype: object

## Summary

Using the pycaret module, we setup the autoML by setting the data to the prepared churn data and target the Churn category of the data set. By using the compare model function, the liner discriminant analysis model provides the best accuracy for the data set. Then a predict model function was used on the last row to make a 2D array. The save model function is used to save the trained model in a pickle file for later use. A test was carried out on the save file by opening the saved pickle with the pickle load function. Using the loaded data, the predict model function is used to make predictions on the new data using the prepared data set. The predictions shows that out of the 5 customers listed, 4 of the customers will churn.